

# Lab 4 - Gaussian Processes

*Arvid Edenheim - arved490*

*2019-08-18*

## Function Setup

```
##### Functions #####

# Squared exponential kernel from lecture examples
# Courtesy of Jose M. Pena, LiU
squaredExponential <- function(x1,x2,sigmaF,l){
  n1 <- length(x1)
  n2 <- length(x2)
  K <- matrix(NA,n1,n2)
  for (i in 1:n2){
    K[,i] <- sigmaF^2 * exp( -0.5 * ( (x1 - x2[i]) / l)^2 )
  }
  return(K)
}

# Prediction function from Rasmussen & Williams
# X = Training inputs
# y = Training targets/outputs
# XStar = Vector of inputs where the posterior distr. is evaluated
# hyperParam = kernel parameters sigma_f and l
# sigmaNoise = Noise sigma_n
# K = Kernel function
posteriorGP <- function(X, y, XStar, sigmaF, l, sigmaNoise, K) {

  # Need to transpose cholesky matrix to attain lower triangular matrix
  L <- t(chol(K(X, X, sigmaF, l) + sigmaNoise^2 * diag(length(X))))
  alpha <- solve(t(L), solve(L, y))

  # Predictive mean
  kStar <- K(X, XStar, sigmaF, l)
  fBar <- t(kStar) %*% alpha

  # Predictive variance
  v <- solve(L, kStar)
  vf <- K(XStar, XStar, sigmaF, l) - t(v) %*% v

  return(list(mean=fBar, variance=diag(vf)))
}

plotGP <- function(XStar, X, y, mean, sd) {
  plot(XStar,
       posterior$mean,
       type='l',
       ylim=c(min(posterior$mean) - 3,max(posterior$mean) + 3))
  points(X, y)
  quantile <- qnorm(0.975)
  lines(XStar, mean + quantile * sd, col='grey')
  lines(XStar, mean - quantile * sd, col='grey')
  legend('topleft', legend = c('Mean', 'Probability band'), fill = c('black', 'grey'))
}
```

## 2.1 - GP Regression implementation

GP model to implement:  $y = f(x) + \epsilon$  with  $\epsilon \sim N(0, \sigma_n^2)$  and  $f \sim GP(0, k(x, x'))$

### 2.1.1

The implementation for the posterior of the GP above was implemented using the algorithm described by Rasmussen & Williams.

```
##### 2.1.1 #####  
# See posteriorGP and squaredExponential above
```

### 2.1.2

The prior was updated with the point  $(x, y) = (0.4, 0.719)$ , using the parameters  $\sigma_f = 1$ ,  $l = 0.3$  and  $\sigma_n = 0.1$ , after which the posterior mean was plotted with 95% probability bands. The generated plot clearly indicates a smaller variance around the observation.

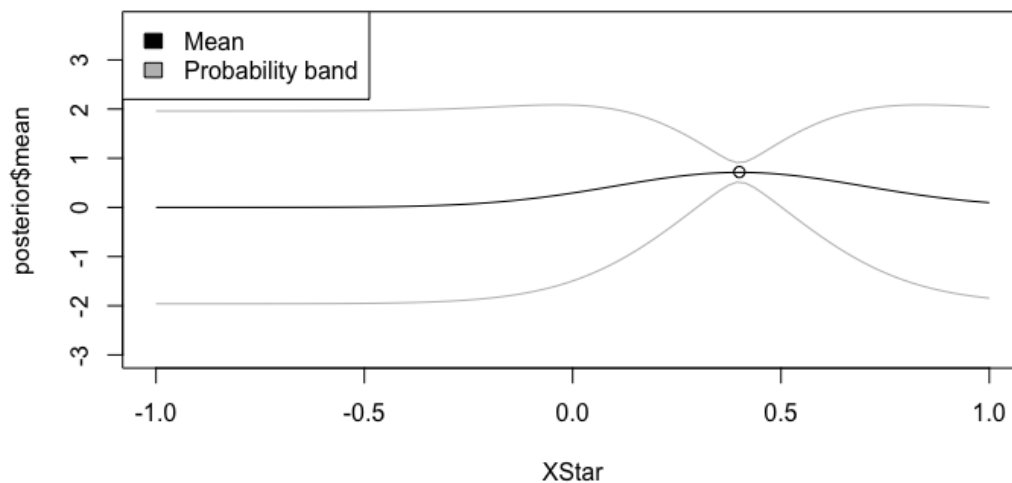


Figure 1: Posterior mean with  $\sigma = 1$  and  $l = 0.3$

```
##### 2.1.2 #####  
sigmaF <- 1  
l <- 0.3  
XStar <- seq(-1, 1, length.out = 100)  
y <- c(0.719)  
X <- c(0.4)  
sigmaN <- 0.1  
  
posterior <- posteriorGP(X, y, XStar, sigmaF, l, sigmaN, squaredExponential)  
plotGP(XStar, X, y, posterior$mean, sqrt(posterior$variance))
```

### 2.1.3

The posterior was updated with yet another observation and plotted. Same conclusion as 2.1.2.

JMP: Updating the posterior after one observation with a new observation gives the same result as updating the prior directly with the two observations.

But that must be less efficient, right? Not even sure how to update the posterior.

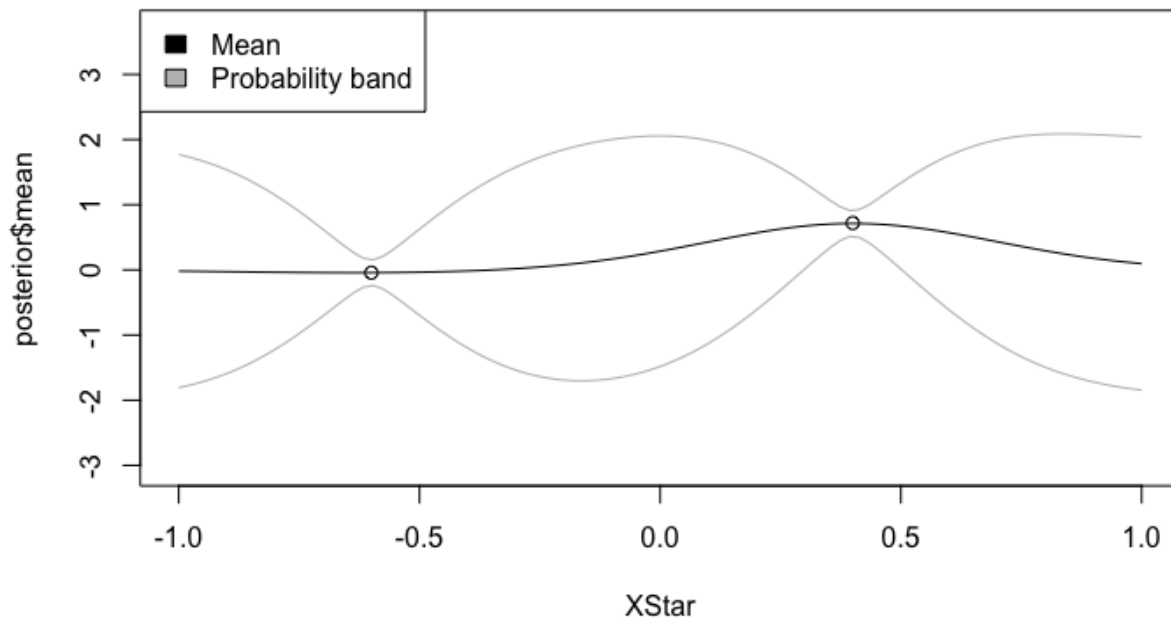


Figure 2: Posterior mean with two observations

```
##### 2.1.3 #####
y <- c(0.719, -0.044)
X <- c(0.4, -0.6)
posterior <- posteriorGP(X, y, XStar, sigmaF, 1, sigmaN, squaredExponential)
plotGP(XStar, X, y, posterior$mean, sqrt(posterior$variance))
```

### 2.1.4

Now with five observations. The posterior starts to resemble a decent regression curve, despite the small number of observations.

x	-1.0	-0.6	-0.2	0.4	0.8
y	0.768	-0.044	-0.940	0.719	-0.664

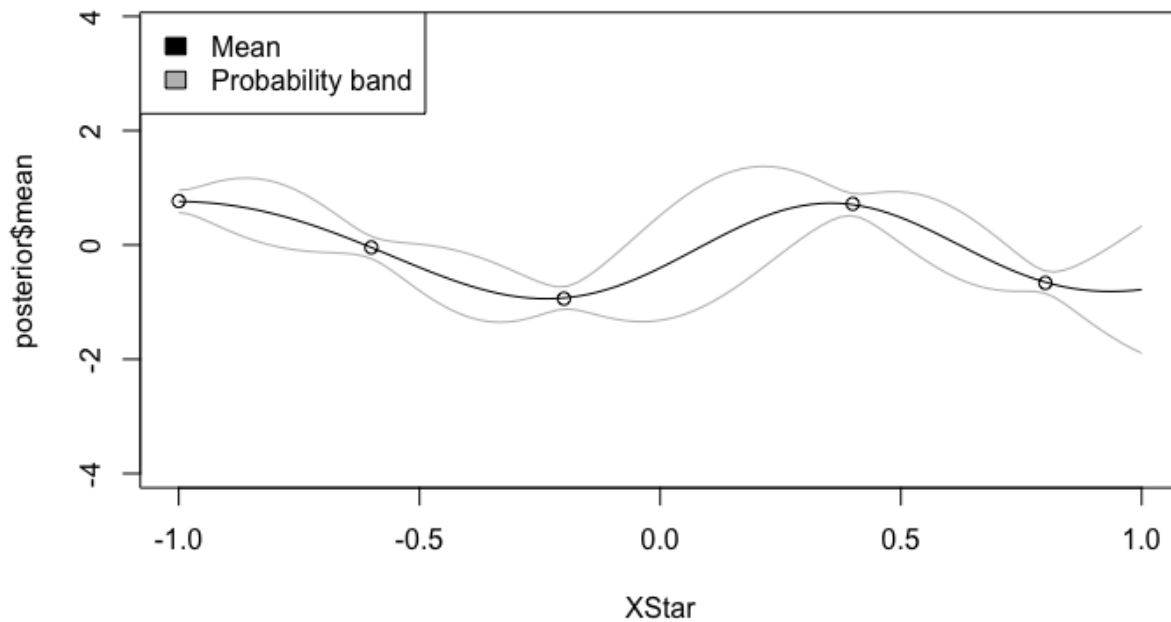


Figure 3: Posterior mean with five observations

```
##### 2.1.4 #####
y <- c(0.719, -0.044, 0.768, -0.940, -0.664)
X <- c(0.4, -0.6, -1, -0.2, 0.8)
posterior <- posteriorGP(X, y, XStar, sigmaF, 1, sigmaN, squaredExponential)
plotGP(XStar, X, y, posterior$mean, sqrt(posterior$variance))
```

### 2.1.5

Same as 2.1.4, but with  $\sigma_f = 1$  and  $l = 1$ .

A higher value for  $l$  gives a wider span in which the observations will affect the posterior (higher covariance), thus decreasing the overall uncertainty/variance in-between observations, making the model smoother, and possibly makes the model more prone to underfit. Looking at the plot,  $l = 1$  seems to result in underfitting the data.

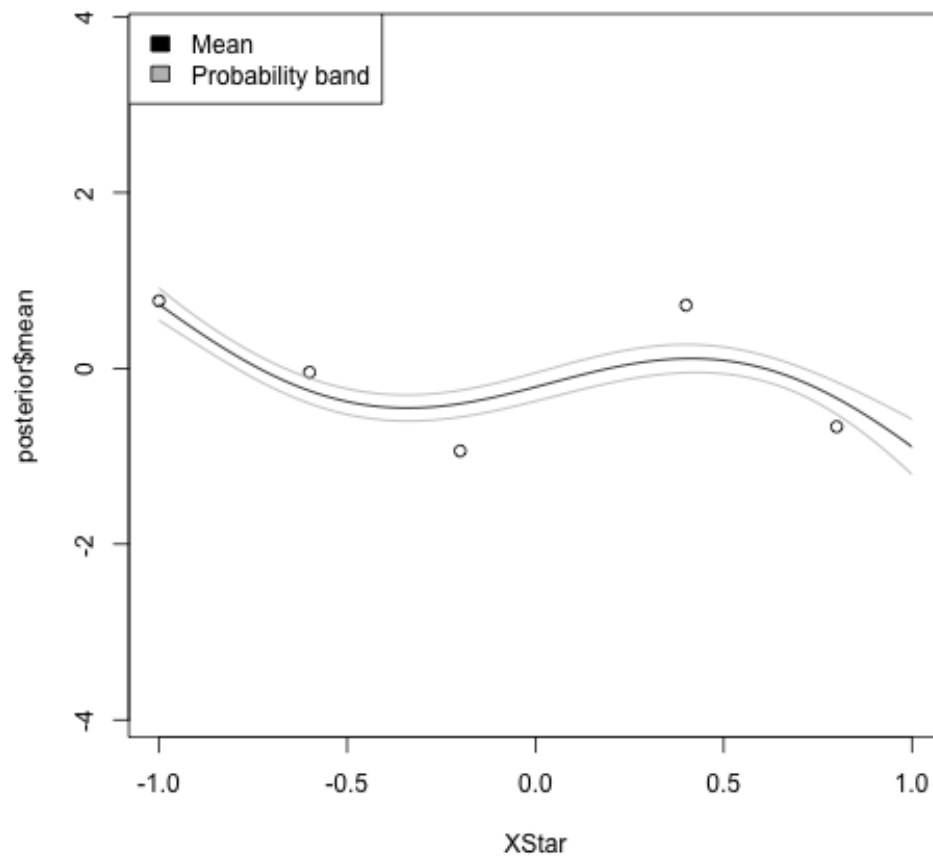


Figure 4: Posterior mean with changed kernel parameters

```
##### 2.1.5 #####  
sigmaF <- 1  
l <- 1
```

## 2.2 - GP Regression with kernlab

### Setup

```
# Setup
temp <- read.csv("https://github.com/STIMaLiU/AdvMLCourse/raw/master/GaussianProcess/Code/TempTullinge.
               header=TRUE,
               sep=";")
time <- 1:2190
day <- rep(1:365, 6)

# Sample (so what I just specified was pointless)
samples <- seq(1, 2190, by=5)
time <- time[samples]
day <- day[samples]
tempSamples <- temp$temp[samples]
```

### 2.2.1

$$K(1, 2) = 0.6065307$$

$$X = (1, 3, 4)^T \text{ \& } X_* = (2, 3, 4)^T$$

$$K(X, X_*) =$$

$$\begin{pmatrix} 0.6065307 & 0.1353353 & 0.0111090 \\ 0.6065307 & 1.0000000 & 0.6065307 \\ 0.1353353 & 0.6065307 & 1.0000000 \end{pmatrix}$$

```
##### 2.2.1 #####
squaredExponentialKernel <- function(ell, sigmaf) {
  squaredExponential <- function(x, y = NULL) {
    n1 <- length(x)
    n2 <- length(y)
    K <- matrix(NA, n1, n2)
    for (i in 1:n2){
      K[,i] <- sigmaf^2 * exp( -0.5 * ( (x - y[i]) / ell)^2 )
    }
    return(K)
  }
  class(squaredExponential) <- 'kernel'
  return(squaredExponential)
}

SEkernel <- squaredExponentialKernel(ell = 1, sigmaf = 1)

SEkernel(1,2)

X <- c(1,3,4)
XStar <- c(2,3,4)

K <- kernelMatrix(kernel = SEkernel,
                  x = X,
                  y = XStar)
```

## 2.2.2

### Model using time as input

$temp = f(time) + \epsilon$  with  $\epsilon \sim N(0, \sigma_n^2)$  and  $f \sim GP(0, k(time, time'))$

Fitted a second degree polynomial due to the characteristics of the data. The generated GP seems to fit the data fairly well.

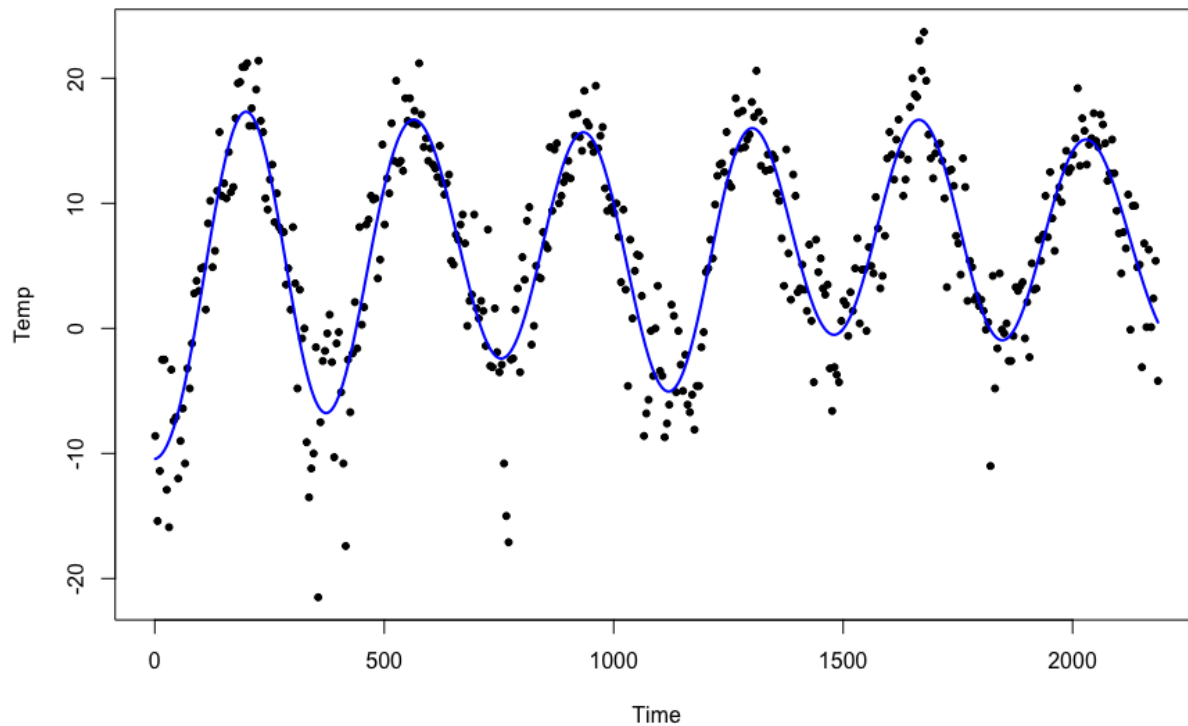


Figure 5: GP regression with time as input

```
##### 2.2.2 #####
plot(temp$temp)

# Fitting a second degree polynomial due to characteristics of the data
fit <- lm(formula=tempSamples ~ time + I(time^2))
sigmaNoise <- sd(fit$residuals)

sigmaF <- 20
l <- 0.2
SEkernel <- squaredExponentialKernel(ell = 1, sigmaf = sigmaF)

GPfit <- gausspr(time, tempSamples, kernel = SEkernel, var = sigmaNoise^2)
meanPred <- predict(GPfit, time) # Predicting the training data. To plot the fit.
```



```
plot(time, tempSamples, pch=20, ylab='Temp', xlab='Time')
lines(time, meanPred, col="blue", lwd = 2)
```

*# As before, increasing l results in underfitting*  
*# Setting a small sigmaF also results in underfitting, otherwise minor effect*

### 2.2.3

I consider the fitted model to be quite good, however underfitting somewhat when considering the extremes.

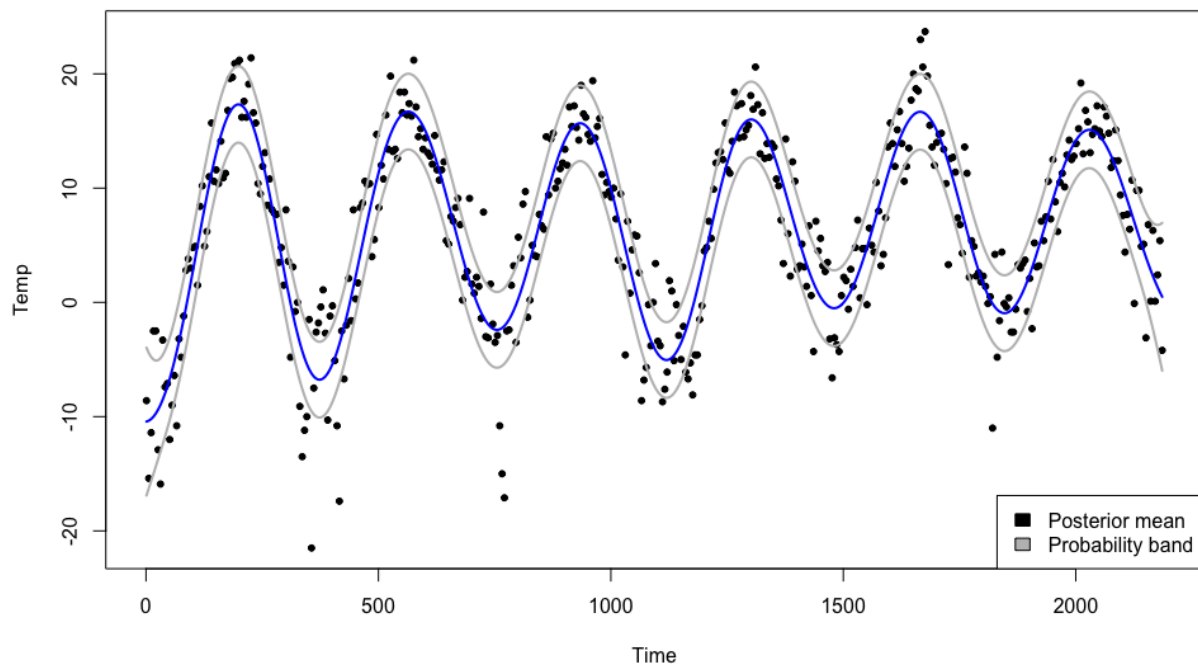


Figure 6: GP with 95% probability bands

```
##### 2.2.3 #####
posterior <- posteriorGP(X = scale(time),
                        y = scale(tempSamples),
                        XStar = scale(time),
                        sigmaF = sigmaF,
                        l = 1,
                        sigmaNoise = sigmaNoise,
                        K = squaredExponential)

quantile <- qnorm(0.025)
lines(time, meanPred - quantile * sqrt(posterior$variance),
      col = "grey", lwd = 2)
lines(time, meanPred + quantile * sqrt(posterior$variance),
```

```
col = "grey", lwd = 2)
legend('bottomright',
      legend = c('Posterior mean', 'Probability band'),
      fill = c('black', 'grey'))
```

## 2.2.4

### Model using day as input

$temp = f(day) + \epsilon$  with  $\epsilon \sim N(0, \sigma_n^2)$  and  $f \sim GP(0, k(day, day'))$

This model also generates a reasonably good model, although not as smooth as in 2.2.2. One possible reason for using this instead would be that it appears to be slightly better at capturing the extremes.

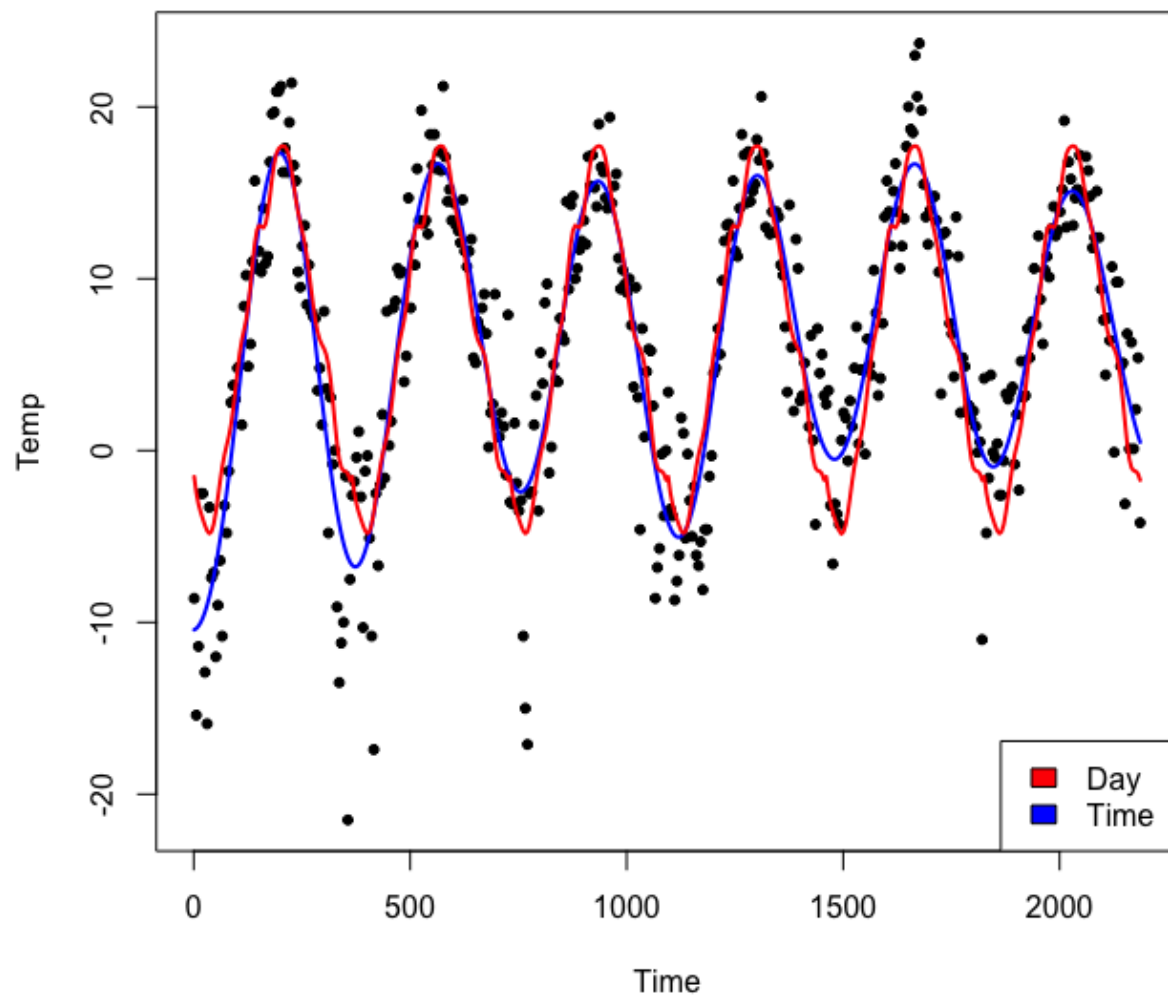


Figure 7: GP regression with day as input

```
##### 2.2.4 #####
fit <- lm(formula=tempSamples ~ day + I(day^2))
sigmaNoise <- sd(fit$residuals)

GPfit <- gausspr(day, tempSamples, kernel = SEkernel, var = sigmaNoise^2)
meanPredDay <- predict(GPfit, day) # Predicting the training data. To plot the fit.

plot(time, tempSamples, pch=20, ylab='Temp', xlab='Time')
lines(time, meanPred, col="blue", lwd = 2)
lines(time, meanPredDay, col="red", lwd = 2)
legend('bottomright', legend = c('Day', 'Time'), fill = c('red', 'blue'))
```

## 2.2.5

The GP with a periodic kernel also appeared to be a good fit. I would prefer this model over the others since it allows more configurability, although increasing the need to select reasonable hyperparameters in order to avoid under/overfitting.

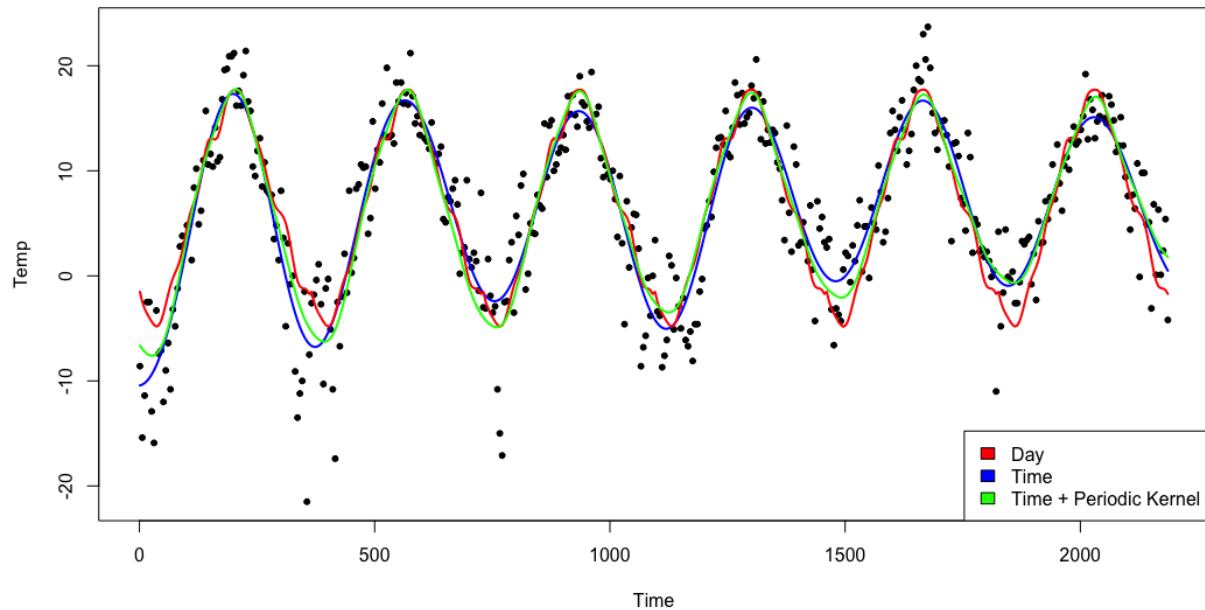


Figure 8: GP with periodic kernel

```
##### 2.2.5 #####

sigmaF <- 20
l1 <- 1
l2 <- 10
d = 365 / sd(time) # Using the sampled interval, nearly the same as sd(1:2190)

PKernel <- function(l1, l2, d, sigmaF) {
  periodicKernel <- function(x, y = NULL) {
    K <- sigmaF^2 * exp(-2 * sin(pi * abs(x - y) / d)^2 / l1^2) * exp(-0.5 * abs(x - y)^2 / l2^2)
    return(K)
  }
  class(periodicKernel) <- 'kernel'
  return(periodicKernel)
}

pk <- PKernel(l1, l2, d, sigmaF)

fit <- lm(formula=tempSamples ~ time + I(time^2))
sigmaNoise <- sd(fit$residuals)
```

```

GPfit <- gausspr(time, tempSamples, kernel = pk, var = sigmaNoise^2)
meanPredPeriodic <- predict(GPfit, time)

plot(time, tempSamples, pch=20, ylab='Temp', xlab='Time')
lines(time, meanPred, col="blue", lwd = 2)
lines(time, meanPredDay, col="red", lwd = 2)
lines(time, meanPredPeriodic, col="green", lwd = 2)
legend('bottomright',
      legend = c('Day', 'Time', 'Time + Periodic Kernel'),
      fill = c('red', 'blue', 'green'))

```

## 2.3 - Classification with kernlab

### Setup

```
# Setup
accuracy <- function(confMatrix) {
  return(sum(diag(confMatrix))/sum(confMatrix))
}

data <- read.csv("https://github.com/STIMALiU/AdvMLCourse/raw/master/GaussianProcess/Code/banknoteFraud.csv",
  header=FALSE,
  sep=",")

names(data) <- c("varWave", "skewWave", "kurtWave", "entropyWave", "fraud")
data[,5] <- as.factor(data[,5])

set.seed(111)
SelectTraining <- sample(1:dim(data)[1], size = 1000, replace = FALSE)
train <- data[SelectTraining,]
test <- data[-SelectTraining,]
```

### 2.3.1 - Using varWave and skewWave to classify

Accuracy = 0.932

Confusion Matrix

	False	True
False	512	24
True	44	420

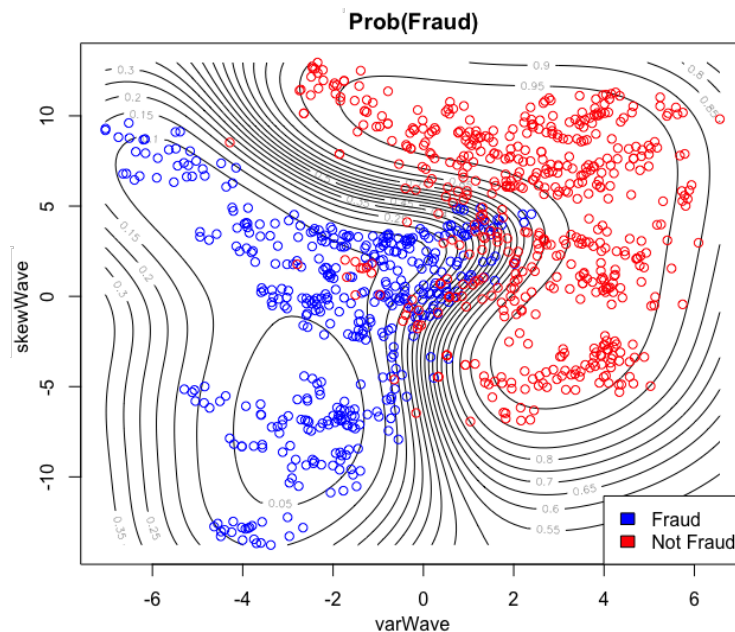


Figure 9: Contour plot of fraud data

```

# Using varWave and skewWave to classify
GPfit <- gausspr(fraud ~ varWave + skewWave, data=train)

# Predict on the training set
GPpred <- predict(GPfit, train[, c(T, T, F, F, F)]) # Doesn't seem to be necessary
confM <- table(GPpred, train[,5]) # confusion matrix
accuracy(confM)

# Class probabilities
probPreds <- predict(GPfit, train[, c(T, T, F, F, F)], type="probabilities")
x1 = seq(min(train$varWave), max(train$varWave), length=100)
x2 = seq(min(train$skewWave), max(train$skewWave), length=100)
gridPoints <- meshgrid(x1, x2)
gridPoints <- cbind(c(gridPoints$x), c(gridPoints$y))
gridPoints <- data.frame(gridPoints)
names(gridPoints) <- names(train)[c(1,2)]
probPreds <- predict(GPfit, gridPoints, type="probabilities")

# Contour plot
contour(x1,
        x2,
        matrix(probPreds[,1], 100, byrow = TRUE),
        20,
        xlab = "varWave",
        ylab = "skewWave",
        main = 'Prob(Fraud)')
points(train$varWave[train$fraud == 1],
        train$skewWave[train$fraud == 1],
        col = 'blue')
points(train$varWave[train$fraud == 0],
        train$skewWave[train$fraud == 0],
        col = 'red')

legend('bottomright', legend = c('Fraud', 'Not Fraud'), fill = c('Blue', 'Red'))

```

### 2.3.2

Accuracy = 0.9354839, which is pretty much the same as the accuracy on the training data. Not great, not terrible.

Confusion Matrix

	False	True
False	191	9
True	15	157

```

##### 2.3.2 #####
# Predict on the test data
GPtest <- predict(GPfit, test[, c(T, T, F, F, F)]) # Doesn't seem to be necessary
confusionMatrix <- table(GPtest, test[,5])
accuracy(confusionMatrix)

```

### 2.3.3

Accuracy on test data = 0.9973118. Great! Just one wrongly predicted out of 372 samples. Overfitting not likely since test data.

Confusion Matrix

	False	True
False	205	0
True	1	166

```
##### 2.3.3 #####
# Use all available covariates
GPfit <- gausspr(fraud ~ varWave + skewWave + kurtWave + entropyWave, data=train)

# Predict on test data
GPtest <- predict(GPfit, test[, c(T, T, T, T, F)]) # Doesn't seem to be necessary
confusionMatrix <- table(GPtest, test[,5])
accuracy(confusionMatrix)
# WOW! High!
```